

Take-Home Assignment

Role: Software Engineer — AI-Native · CarDekho Group

The Brief

A car research platform has a dataset of cars — makes, models, variants, prices, specs, mileage, safety ratings, and user reviews. Buyers come to the platform confused: there are too many options and no easy way to figure out which car is right for them.

Build something that helps a car buyer go from “I don’t know what to buy” to “I’m confident about my shortlist.”

That’s the entire brief. It’s deliberately vague. What you choose to build, what you cut, and how you scope it — that’s the test.

Hard Constraints

1. **Working web app.** It must run — either deployed somewhere (Vercel, Railway, Netlify, whatever) or with a single-command local setup (`docker compose up` , `npm run dev` , etc.). If we can’t run it in under 2 minutes, it doesn’t count.
 2. **Full-stack.** There must be a frontend a human can interact with and a backend that does something non-trivial (persistence, computation, aggregation — not just a static page).
 3. **Screen recording is mandatory.** Record your entire build process — terminal, editor, browser, everything. We want to see how you use your tools (Claude Code, Cursor, Copilot, vanilla vim, whatever). Unedited, real-time or lightly fast-forwarded. Upload to Loom, Google Drive, or YouTube (unlisted). **This is the most important deliverable.**
 4. **Time-box: 2-3 hours of actual work.** Don’t spend a weekend on this. We’re evaluating how much you ship in a constrained window, not how polished you can make it with unlimited time.
-

Deliverables

Send us:

- **Link to the screen recording** (most important)
 - **GitHub repo** (public or invite us)
 - **Live URL** or clear run instructions
 - **A README** that answers:
 - What did you build and why? What did you deliberately cut?
 - What's your tech stack and why did you pick it?
 - What did you delegate to AI tools vs. do manually? Where did the tools help most? Where did they get in the way?
 - If you had another 4 hours, what would you add?
-

What We're Evaluating

Transparent about this — no gotchas.

Product decisions (35%) Did you make smart scoping calls? Did you build the highest-value thing first? Does the thing you shipped actually help a confused car buyer make progress? We'd rather see a tight, opinionated MVP than a half-finished kitchen sink.

Agentic tool usage (30%) This is why the screen recording matters. We want to see: How do you prompt? How do you break down work for the AI? Do you blindly accept output or do you review and course-correct? How do you recover when the tool gives you garbage? Do you know when to stop using the tool and just write code?

Code quality & architecture (20%) Is the code readable and sensible? Are there obvious patterns or is it spaghetti? We're not looking for perfection — we're looking for evidence that you understand what good code looks like even when AI wrote most of it. Bonus if you set up linting, types, or tests — but only if it was fast, not ceremonial.

Execution speed & shipping instinct (15%) How much did you actually get done in 2-3 hours? Is it deployed? Does it work end-to-end? We value "it works and it's live" over "it's architecturally beautiful but I ran out of time."

What We Don't Care About

- Which language or framework you used
 - Whether the UI is pixel-perfect
 - Whether you wrote tests for everything
 - How many years of React experience you have
 - Whether you know anything about the automotive industry
-

Anti-Cheating Note

We know AI tools can build this entire thing. That's the point — we *want* you to use them. The screen recording exists not to catch cheating but to see *how you think and work*. There is no way to cheat on this assignment. There is only your process, and it either impresses us or it doesn't.

Timeline

You have **48 hours** from receiving this assignment to submit your deliverables. The actual build should take 2-3 hours — the rest is buffer for your schedule.